

1. What is an operating system, and what are its primary functions?

An Operating System (OS) is system software that acts as an intermediary between users and computer hardware.

Primary Functions:

- Process Management – Creates, schedules, and terminates processes.
 - Memory Management – Allocates and manages memory efficiently.
 - File System Management – Manages files and directories.
 - I/O Device Management – Controls peripheral devices.
 - Security & Access Control – Protects against unauthorized access.
-

2. Explain the difference between a process and a thread.

- Process: Independent execution unit with its own memory space.
- Thread: A lightweight unit within a process that shares memory and resources.

Comparison:

Feature	Process	Thread
Memory	Has its own memory space	Shares memory within a process
Context Switching	Slow	Fast
Communication	Uses IPC (Inter-Process Communication)	Direct memory access
Creation	Expensive	Cheaper

3. What is virtual memory, and how does it work?

Virtual memory is a technique that provides the illusion of more memory than physically available by using disk space (swap space).

- Paging – Divides memory into fixed-size pages and swaps them between RAM and disk.
 - Segmentation – Divides memory logically into variable-sized segments.
-

4. Describe the difference between multiprogramming, multitasking, and multiprocessing.

Feature	Multiprogramming	Multitasking	Multiprocessing
Definition	Multiple programs in memory	Multiple tasks running	Multiple CPUs
Execution	Only one process at a time	Rapid switching of tasks	True parallel execution
Goal	Maximize CPU utilization	Increase responsiveness	Improve performance

5. What is a file system, and what are its components?

A file system organizes and manages data storage.

Components:

- Files – Data storage units.
- Directories – Organizational structures.
- File Allocation Table (FAT) / Inode Table – Stores metadata.
- File Access Permissions – Controls access.

Types of File Systems: FAT32, NTFS, EXT4, HFS+

6. What is a deadlock, and how can it be prevented?

A deadlock occurs when processes hold resources and wait indefinitely for each other.

Deadlock Prevention Strategies:

- Avoid Hold and Wait – Request all resources at once.
 - Circular Wait Prevention – Impose an ordering rule.
 - Deadlock Detection & Recovery – Detect cycles and terminate processes.
-

7. Explain the difference between a kernel and a shell.

- Kernel: The core of the OS that interacts with hardware.
- Shell: The interface that users interact with to execute commands.

Types of Shells:

- Command-Line Shell (e.g., Bash, PowerShell)
- Graphical Shell (e.g., Windows Explorer)

8. What is CPU scheduling, and why is it important?

CPU Scheduling determines which process gets CPU time to optimize performance.

Scheduling Algorithms:

- FCFS (First-Come-First-Serve)
- SJF (Shortest Job First)
- Round Robin
- Priority Scheduling

9. How does a system call work?

A system call allows user programs to request services from the OS kernel.

Examples:

- `open()`, `read()`, `write()` – File operations
- `fork()`, `exec()` – Process creation
- `kill()` – Process termination

10. What is the purpose of device drivers in an operating system?

Device drivers allow the OS to communicate with hardware (printers, keyboards, etc.).

- Kernel-mode drivers – Run at a low level, directly interacting with hardware.
- User-mode drivers – Operate in user space for safer execution.

11. Explain the role of the page table in virtual memory management.

A page table maps virtual addresses to physical memory locations.

- Page Table Entry (PTE): Stores frame number, valid bit, access rights.
 - Translation Lookaside Buffer (TLB): Caches page table entries for faster access.
-

12. What is thrashing, and how can it be avoided?

Thrashing occurs when excessive paging degrades system performance.

Solutions:

- Increase RAM.
 - Implement Working Set Model (only keep frequently used pages in memory).
 - Use Page Fault Frequency control.
-

13. Describe the concept of a semaphore and its use in synchronization.

A semaphore is a signaling mechanism used to manage process synchronization.

Types:

- Binary Semaphore (Mutex): Values 0 or 1, used for mutual exclusion.
 - Counting Semaphore: Allows multiple processes to access resources.
-

14. How does an operating system handle process synchronization?

Using synchronization mechanisms like:

- Semaphores – Prevent race conditions.
 - Monitors – Encapsulate shared resources.
 - Locks & Mutexes – Ensure atomicity.
-

15. What is the purpose of an interrupt in operating systems?

Interrupts temporarily halt CPU execution to handle urgent tasks like I/O operations.

Types:

- Hardware Interrupts – Triggered by external devices.
 - Software Interrupts – Generated by programs (e.g., system calls).
-

16. Explain the concept of a file descriptor.

A file descriptor is an integer that represents an open file in Unix-like systems.

File Descriptor Standard Stream

0	Standard Input (stdin)
1	Standard Output (stdout)
2	Standard Error (stderr)

17. How does a system recover from a system crash?

- File System Journaling – Logs changes before applying them.
 - Checkpointing – Saves system state periodically.
 - Process Rollback – Restores the last stable state.
-

18. Describe the difference between a monolithic kernel and a microkernel.

Feature	Monolithic Kernel	Microkernel
Structure	Single large program	Minimal core, user-space modules
Performance	Fast	Slower due to IPC overhead
Stability	Less stable	More stable

19. What is demand paging, and how does it improve memory management efficiency?

Demand paging loads only the required memory pages, reducing memory usage.

Advantages:

- Saves RAM.
 - Reduces I/O overhead.
 - Faster process loading.
-

20. Explain the concept of a race condition and how it can be prevented.

A race condition occurs when multiple processes access shared data simultaneously, leading to inconsistent results.

Prevention Methods:

- Locks & Mutexes – Ensure exclusive access.
- Atomic Operations – Prevent interruptions.
- Critical Sections – Restrict concurrent access.

21. Explain the difference between preemptive and non-preemptive scheduling.

- Preemptive Scheduling: The OS can interrupt and switch a process before it completes (e.g., Round Robin, Priority Scheduling).
 - Non-Preemptive Scheduling: Once a process starts execution, it cannot be interrupted until it finishes (e.g., FCFS, SJF).
-

22. What is round-robin scheduling, and how does it work?

Round-Robin (RR) is a preemptive scheduling algorithm where each process gets a fixed time quantum for execution.

- If the process is not completed within the time quantum, it is moved to the end of the queue.
 - It ensures fair CPU time allocation and is used in time-sharing systems.
-

23. Describe the priority scheduling algorithm. How is priority assigned to processes?

- Each process is assigned a priority value; the scheduler selects the highest priority process.
- Two Types:
 - Preemptive Priority Scheduling: A higher-priority process interrupts a lower-priority one.
 - Non-Preemptive Priority Scheduling: A higher-priority process waits until the current process finishes.

Priority Assignment Methods:

- Static: Assigned at process creation.
 - Dynamic: Adjusted based on execution time or resource needs.
-

24. What is the shortest job next (SJN) scheduling algorithm, and when is it used?

- SJN (or SJF – Shortest Job First) selects the process with the shortest CPU burst time.
 - Used in: Batch systems where process execution times are known in advance.
 - Problem: Can lead to starvation for longer jobs.
-

25. Explain the concept of multilevel queue scheduling.

- Divides processes into separate queues based on priority/type (e.g., system processes, interactive, batch).
 - Each queue can have its own scheduling algorithm (e.g., RR for interactive, FCFS for batch).
 - Ensures better process categorization and execution control.
-

26. What is a process control block (PCB), and what information does it contain?

A PCB stores process-related information required by the OS.

Contains:

- Process ID (PID)
 - Process state (Ready, Running, Waiting)
 - Program counter (Next instruction)
 - CPU registers & scheduling information
 - Memory management info
-

27. Describe the process state diagram and transitions between process states.

State	Transition	Description
New	Created → Ready	Process is created.
Ready	CPU assigned → Running	Waiting for CPU time.
Running	Time expired → Ready OR I/O needed → Waiting	Currently executing.
Waiting	I/O complete → Ready	Waiting for I/O or resource.
Terminated	Process completed	Process exits.

28. How does a process communicate with another process in an operating system?

Processes communicate using Inter-Process Communication (IPC) mechanisms:

- Pipes – Data transfer between related processes.
 - Shared Memory – Common memory space for direct communication.
 - Message Queues – Message-passing between processes.
 - Sockets – Communication between processes across networks.
-

29. What is process synchronization, and why is it important?

- Ensures that multiple processes execute correctly when accessing shared resources.
 - Avoids issues like race conditions, deadlocks, and inconsistency.
 - Uses semaphores, mutexes, monitors, and locks.
-

30. Explain the concept of a zombie process and how it is created.

A zombie process occurs when:

- A child process terminates but its parent does not call `wait()` to collect its exit status.
- The child remains in the process table as a defunct (zombie) process.

Prevention: Use `wait()` or `waitpid()` in the parent process.

31. Describe the difference between internal fragmentation and external fragmentation.

- Internal Fragmentation: Wasted space inside allocated memory blocks (e.g., a 100-byte block used for 80 bytes).
- External Fragmentation: Free memory exists but is scattered in non-contiguous blocks.

Solution: Compaction or paging.

32. What is demand paging, and how does it improve memory management efficiency?

Demand paging loads pages into memory only when required, reducing memory usage.

- Uses Page Fault Handling to fetch missing pages.
 - Advantages: Reduces memory wastage, speeds up execution.
-

33. Explain the role of the page table in virtual memory management.

A page table maps virtual addresses to physical memory addresses.

Types:

- Single-level paging – Simple but slow.
 - Multi-level paging – Reduces page table size.
 - Inverted page table – Uses a hash table for efficient lookups.
-

34. How does a memory management unit (MMU) work?

The MMU:

- Converts virtual addresses to physical addresses.
 - Uses TLB (Translation Lookaside Buffer) for fast lookups.
 - Supports paging and segmentation for memory protection.
-

35. What is thrashing, and how can it be avoided in virtual memory systems?

- Thrashing occurs when too many page faults degrade performance.
 - Prevented by:
 - Increasing RAM.
 - Adjusting the working set size.
 - Using page replacement algorithms like LRU or FIFO.
-

36. What is a system call, and how does it facilitate communication between user programs and the OS?

A system call is a request to the OS for privileged operations like:

- File management (open(), read(), write()).
 - Process control (fork(), exec(), exit()).
 - Memory management (mmap(), brk()).
-

37. How does an operating system handle I/O operations?

- Uses buffering, caching, and spooling.
- Employs interrupt-driven and DMA (Direct Memory Access) techniques.

38. Explain the concept of a race condition and how it can be prevented.

- A race condition occurs when multiple threads/processes access shared data simultaneously.
- Prevented using synchronization mechanisms (semaphores, mutexes, locks).

39. Describe the role of device drivers in an operating system.

- Acts as a translator between the OS and hardware.
- Handles I/O operations, interrupts, and device communication.

40. What is an orphan process? How does an OS handle orphan processes?

- A child process becomes orphaned if its parent terminates before it.
- Solution: The OS reassigns it to init (PID 1) to prevent resource leaks.

41. How does the fork() system call work in creating a new process in Unix-like operating systems?

- fork() creates a new child process that is an exact copy of the parent.
- The child receives a unique Process ID (PID).
- Both parent and child continue execution.

42. How does process termination occur in Unix-like operating systems?

- Normal exit (exit()).
- Signal-based termination (kill(pid, signal)).
- Parent termination – If the parent dies, the child becomes orphaned.

43. What is the role of the long-term scheduler in process scheduling?

- Controls the degree of multiprogramming by selecting processes from disk to be loaded into memory.
-

44. How does the short-term scheduler differ from the long-term scheduler?

Feature	Long-Term Scheduler	Short-Term Scheduler
Purpose	Decides which processes to load	Selects which process to run next
Execution Frequency	Infrequent	Frequent
Affects	Multiprogramming	CPU scheduling

45. Describe how a parent process can wait for a child process to finish execution.

A parent process can use the `wait()` or `waitpid()` system calls to wait for its child process to complete.

- `wait()`: Blocks the parent until any child process terminates.
 - `waitpid(pid, &status, options)`: Allows waiting for a specific child process.
 - Prevents zombie processes by releasing resources used by the child.
-

46. What is the significance of the exit status of a child process in the `wait()` system call?

- The exit status informs the parent process why the child terminated.
 - It is stored in an integer variable passed to `wait()` or `waitpid()`.
 - Example:
 - `exit(0)`: Normal termination.
 - `exit(1)`: Error occurred.
 - Termination by signal (e.g., segmentation fault).
-

47. How can a parent process terminate a child process in Unix-like operating systems?

- Using the `kill(pid, signal)` system call:
 - `kill(child_pid, SIGTERM)`: Graceful termination.
 - `kill(child_pid, SIGKILL)`: Forceful termination.
 - The parent can also terminate all child processes when it exits.
-

48. Explain the difference between a process group and a session in Unix-like operating systems.

- Process Group: A collection of related processes that can receive the same signals together.
 - Session: A set of process groups managed by a controlling terminal.
 - Example:
 - A shell runs a process (session leader).
 - A session may contain multiple process groups (foreground and background jobs).
-

49. Describe how the `exec()` family of functions is used to replace the current process image with a new one.

- The `exec()` functions replace the current process image with a new program.
- Common variants:
 - `execl()`, `execp()`, `execv()`, etc.
- Example Usage:

c

CopyEdit

```
execl("/bin/l", "l", "-l", NULL);
```

- This replaces the running process with `/bin/l`.
-

50. What is the purpose of the `waitpid()` system call in process management? How does it differ from `wait()`?

- `waitpid()` allows a parent to wait for a specific child process, while `wait()` waits for any child.
- Syntax:

c

CopyEdit

```
pid_t waitpid(pid, &status, options);
```

- Differences:
 - `waitpid(pid, NULL, 0)`: Waits for the specific process.
 - `waitpid(-1, NULL, 0)`: Acts like `wait()`, waits for any child.
-

51. How does process termination occur in Unix-like operating systems?

A process terminates when:

1. Normal exit (`exit()`).
2. Signal-based termination (`kill(pid, SIGTERM)`).
3. Fatal error (e.g., segmentation fault).
4. Parent process termination (orphaned processes get adopted by `init`).

52. What is the role of the long-term scheduler in the process scheduling hierarchy? How does it influence the degree of multiprogramming in an operating system?

- Long-Term Scheduler (Job Scheduler) controls the number of processes in memory.
- Determines the degree of multiprogramming (how many processes are in RAM).
- Helps balance CPU-bound and I/O-bound processes for efficient execution.

53. How does the short-term scheduler differ from the long-term and medium-term schedulers in terms of frequency of execution and the scope of its decisions?

Feature	Long-Term Scheduler	Short-Term Scheduler	Medium-Term Scheduler
Purpose	Loads processes into memory	Schedules processes for CPU execution	Swaps processes in/out of memory
Execution Frequency	Infrequent	Very Frequent	Occasional
Affects	Degree of multiprogramming	CPU scheduling	Memory management

54. Describe a scenario where the medium-term scheduler would be invoked and explain how it helps manage system resources more efficiently.

- The medium-term scheduler temporarily removes processes from RAM (swapping).
- Scenario:
 - If too many processes cause thrashing, the scheduler swaps out some processes to free memory.
 - Once resources are available, swapped-out processes are brought back into memory.
- Goal: Improve system performance by optimizing CPU and memory usage.

